

# Java 8 Parallel Streams

Ethan Williams

February 20, 2018

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Splitterator</b>	<b>2</b>
2.1	Basics . . . . .	2
2.2	Implementation . . . . .	3
<b>3</b>	<b>Collectors</b>	<b>5</b>
<b>4</b>	<b>Practical Considerations when Using Parallel Streams</b>	<b>5</b>

## Prerequisites

### Technical

This document was written for Java developers who have an interest in using concurrency (parallelism) in streams. Developers in other languages such as C# may also find the topics useful with the understanding that syntax, implementation, and functionality may/will differ.

Additionally, the author assumes knowledge of serial streams in Java (the default stream implementation). Knowledge of concurrency tools/practices in Java is encouraged but not necessary.

### Language

Some of the language used in this paper may be ambiguous/confusing, the author's intended meanings for these terms are below:

- A stream instantiated with `.stream()` only is referred to as a serial stream
- A stream instantiated with `.parallel().stream()` or `.parallelStream()` is referred to as a parallel stream

## 1 Introduction

Parallel streams were introduced into Java 8 alongside serial streams so that developers could utilize concurrency to make stream processing faster.

## 2 Spliterator

The Spliterator is the backbone of parallel streams, allowing the program to split a collection apart and iterate through it. If a collection object does not have a `spliterator()` method returning a Spliterator object, then Java is not able to process the collection with a parallel stream at all.

### 2.1 Basics

A Spliterator is as an object which is represents the elements in a given range of the backing collection. At the beginning of a parallel stream there is one Spliterator that represents the entire collection. Calling the method `trySplit()` will return a new Spliterator that represents the first half of untraversed elements while modifying the instance to represent the second half of untraversed elements. Figure 1 illustrates a call to `trySplit()` on a Spliterator representing the entire collection.

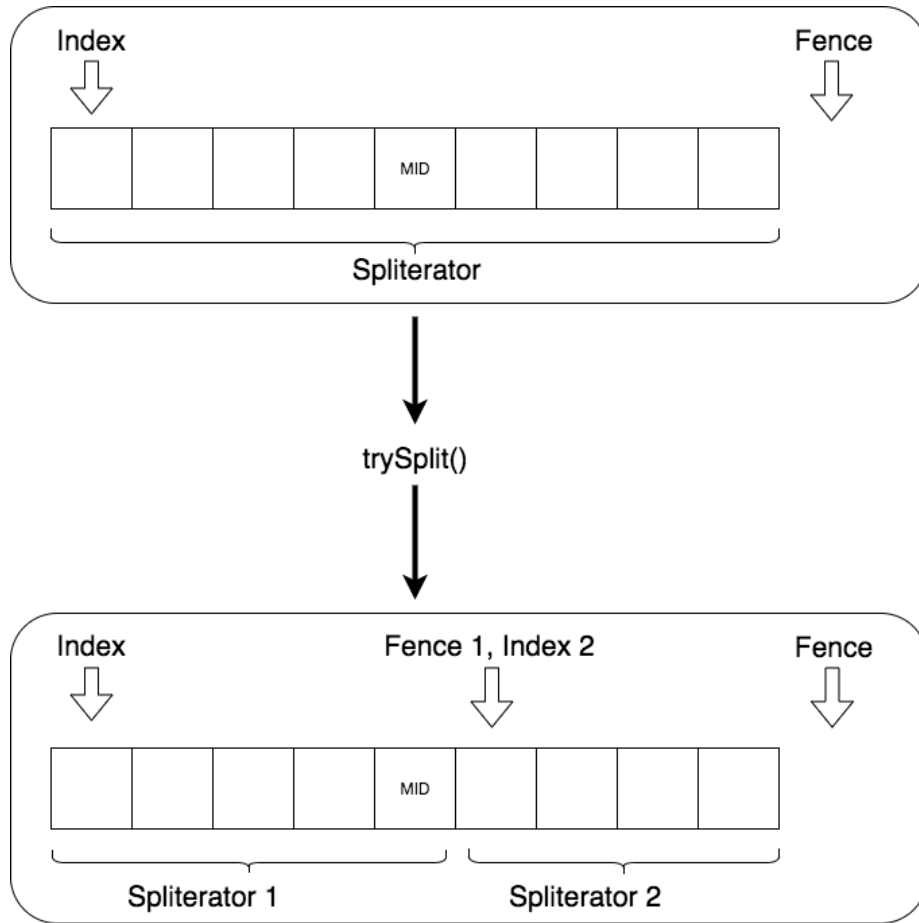


Figure 1: A Spliterator before and after splitting

## 2.2 Implementation

*Spliterator implementations differ based on the collection so this section focuses on a simplified (easier to read) implementation of the Spliterator for Java's `ArrayList`.*

`Spliterator` has 3 instance fields: a backing collection, an index, and a fence. The backing collection for every `Spliterator` is the entire collection it is based on. The index is the current cursor position of the `Spliterator`, which is analogous to the cursor position in an `Iterator` object. The fence is the index of the last element the `Spliterator` is responsible for plus one. The primary methods are `tryEachRemaining()`, `tryNext()`, & `trySplit()`.

Although the cursor is similar to `Iterator`'s, traversal through a `Spliterator` using `tryEachRemaining()` and `tryNext()` is different. `tryEachRemaining()` in Figure 2 takes a `Consumer` object which is the operation to be carried out

on each element of the collection. A typical implementation will simply iterate through all elements and call the `accept()` method on the consumer with the element as its only parameter. `tryNext()` in Figure 3 is similar although the operation is only attempted on element at the current cursor position. If that cursor position is past the fence of the `Splitterator`, then the method returns false, otherwise it returns true.

```
public void forEachRemaining(Consumer<? super E> action) {
    int i;
    if ((i = index) >= 0 && (index = fence) <= a.length) {
        for (; i < hi; ++i) action.accept((E) list.elementData[i]);
    }
    throw new ConcurrentModificationException();
}
```

Figure 2: Implementation of `forEachRemaining()`

```
public boolean tryAdvance(Consumer<? super E> action) {
    int hi = getFence(), i = index;
    if (i < hi) {
        index = i + 1;
        action.accept((E) list.elementData[i]);
        return true;
    }
    return false;
}
```

Figure 3: Implementation of `tryAdvance()`

`Splitterator`'s primary functionality is encapsulated within the `trySplit()` method in Figure 4

```
public Splitterator<E> trySplit() {
    int lo = index, mid = (lo + fence) >>> 1;
    return (lo >= mid) ? null : new Splitterator<E>(list, lo, index = mid);
}
```

Figure 4: Implementation of `trySplit()`

### **3 Collectors**

### **4 Practical Considerations when Using Parallel Streams**